



Evaluation of Relational Operations

Chapter 14, Part A (Joins)



Relational Operations

- ❖ We will consider how to implement:
 - Selection (σ) Selects a subset of rows from relation.
 - Projection (π) Deletes unwanted columns from relation.
 - Join (\bowtie) Allows us to combine two relations.
 - Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
 - Union (\cup) Tuples in reln. 1 and in reln. 2.
 - Aggregation (SUM, MIN, etc.) and GROUP BY



Schema for Examples

Sailors (sid: integer, sname: string, rating: integer, age: real)
Reserves (sid: integer, bid: integer, day: dates, rname: string)


❖ Reserves:

- Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.

❖ Sailors:


- Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

Equality Joins With One Join Column



```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```


- ❖ In algebra: $R \bowtie S$. Common! Must be carefully optimized. $R \times S$ is large; so, $R \times S$ followed by a selection is inefficient.
- ❖ Assume: M pages for R , p_R tuples per page, N pages for S , p_S tuples per page.
 - In our examples, R is Reserves and S is Sailors.
- ❖ *Cost metric*: # of I/Os. We will ignore output costs.



Join Algorithms to Consider


- ❖ Nested loop join
- ❖ Sort-merge join
- ❖ Hash join
- ❖ Index nested loop join

Simple Nested Loops Join




```
foreach tuple r in R do
    foreach tuple s in S do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```


- ❖ Tuple-oriented NLJ: For each tuple in the *outer* relation R, we scan the entire *inner* relation S.
 - Cost: $M + p_R * M * N = 1000 + 100 * 1000 * 500$ I/Os.
- ❖ Page-oriented NLJ: For each *page* of R, get each *page* of S, and write out matching pairs of tuples $\langle r, s \rangle$, where r is in R-page and S is in S-page.
 - Cost: $M + M * N = 1000 + 1000 * 500$
 - If smaller relation (S) is outer, cost = $500 + 500 * 1000$



Block Nested Loops Join


- ❖ Use one page as an input buffer for scanning the inner S, one page as the output buffer, and use all remaining pages to hold “block” of outer R.
 - For each matching tuple r in R-block, s in S-page, add $\langle r, s \rangle$ to result. Then read next R-block, scan S, etc.






Block Nested Loops Join


- ❖ R is scanned once, cost = M pages
- ❖ Each block has size B-2 (B is # of buffer pages), so need to read S a total of $\text{ceiling}(M/[B-2])$ times
- ❖ Total cost: $M + N * \text{ceiling}(M/[B-2])$






Examples of Block Nested Loops

- ❖ **Cost: Scan of outer + #outer blocks * scan of inner**
 - $\#outer\ blocks = \lceil \# of\ pages\ of\ outer / blocksize \rceil$
- ❖ With Reserves (R) as outer, and 100 pages of R:
 - Cost of scanning R is 1000 I/Os; a total of 10 *blocks*.
 - Per block of R, we scan Sailors (S); 10*500 I/Os.
 - If space for just 90 pages of R, we would scan S 12 times.
- ❖ With 100-page block of Sailors as outer:
 - Cost of scanning S is 500 I/Os; a total of 5 blocks.
 - Per block of S, we scan Reserves; 5*1000 I/Os.



Sort-Merge Join $(R \bowtie_{i=j} S)$

- ❖ Sort R and S on the join column, then scan them to do a “merge” (on join col.), and output result tuples.
 - Advance scan of R until current R-tuple \geq current S tuple, then advance scan of S until current S-tuple \geq current R tuple; do this until current R tuple = current S tuple.
 - At this point, all R tuples with same value in R_i (*current R group*) and all S tuples with same value in S_j (*current S group*) match; output $\langle r, s \rangle$ for all pairs of such tuples.
 - Then resume scanning R and S.
- ❖ R is scanned once; each S group is scanned once per matching R tuple. (Multiple scans of an S group are likely to find needed pages in buffer.)



Example of Sort-Merge Join

| <u>sid</u> | <u>bid</u> | <u>day</u> | rname |
|------------|------------|------------|--------|
| 28 | 103 | 12/4/96 | guppy |
| 28 | 103 | 11/3/96 | yuppy |
| 31 | 101 | 10/10/96 | dustin |
| 31 | 102 | 10/12/96 | lubber |
| 31 | 101 | 10/11/96 | lubber |
| 58 | 103 | 11/12/96 | dustin |


| <u>sid</u> | sname | rating | age |
|------------|--------|--------|------|
| 22 | dustin | 7 | 45.0 |
| 28 | yuppy | 9 | 35.0 |
| 31 | lubber | 8 | 55.5 |
| 44 | guppy | 5 | 35.0 |
| 58 | rusty | 10 | 35.0 |



Example of Sort-Merge Join

| <u>sid</u> | sname | rating | age | <u>sid</u> | <u>bid</u> | <u>day</u> | rname |
|------------|--------|--------|------|------------|------------|------------|--------|
| 22 | dustin | 7 | 45.0 | 28 | 103 | 12/4/96 | guppy |
| 28 | yuppy | 9 | 35.0 | 28 | 103 | 11/3/96 | yuppy |
| 31 | lubber | 8 | 55.5 | 31 | 101 | 10/10/96 | dustin |
| 44 | guppy | 5 | 35.0 | 31 | 102 | 10/12/96 | lubber |
| 58 | rusty | 10 | 35.0 | 31 | 101 | 10/11/96 | lubber |
| | | | | 58 | 103 | 11/12/96 | dustin |

- ❖ Cost: $M \log M + N \log N + (M+N)$ [this is incorrect!]
 - The cost of scanning, $M+N$, could be $M*N$ (very unlikely!)




Cost of Sort-Merge Join

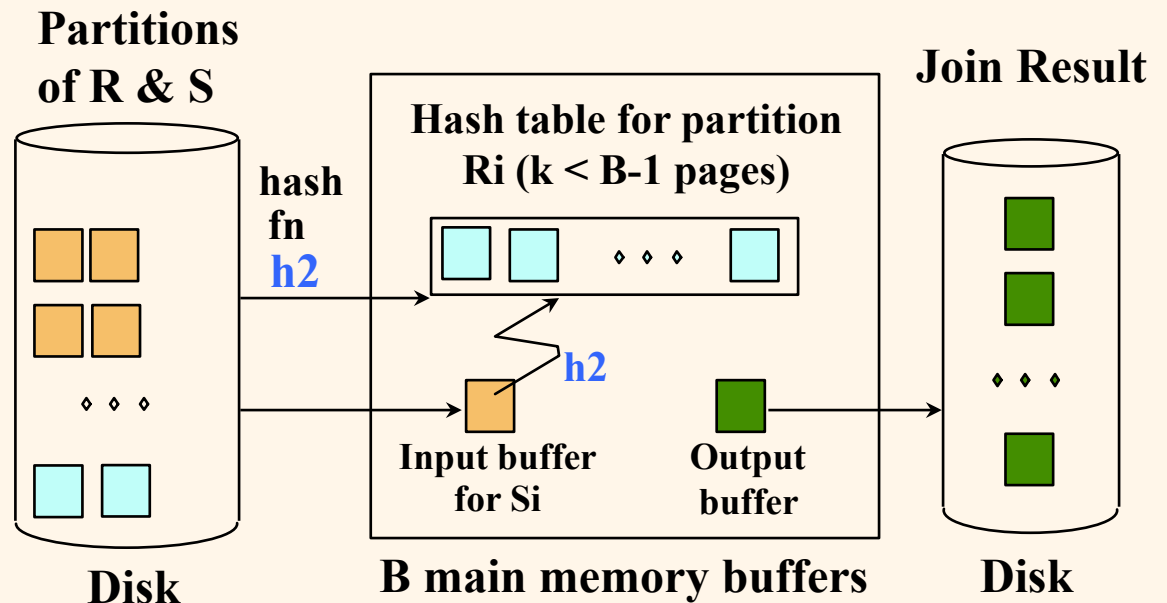
- ❖ Start by sorting both R and S on the join attribute:
 - Assumption: $M \leq B^2$, $N \leq B^2$
 - Cost: $4M + 4N$
- ❖ Read both relations in sorted order, match tuples
 - Typical cost: $M + N$ (but can be as high as $M \cdot N$)
- ❖ Difficulty: many tuples in R may match many in S
 - If at least one set of tuples fits in memory, we are OK
 - Otherwise need nested loop, higher cost
- ❖ Total cost: $5M + 5N$


Hash-Join

❖ Partition both relations using hash fn **h**: R tuples in partition *i* will only match *S* tuples in partition *i*.




❖ Read in a partition of R, hash it using **h2** (\neq **h**!). Scan matching partition of S, search for matches.






Observations on Hash-Join

- ❖ #partitions $k < B-1$ (why?), and $B-2 >$ size of largest partition to be held in memory. Assuming uniformly sized partitions, and maximizing k , we get:
 - $k = B-1$, and $M/(B-1) < B-2$, i.e., B must be $> \sqrt{M}$
- ❖ If we build an in-memory hash table to speed up the matching of tuples, a little more memory is needed.
- ❖ If the hash function does not partition uniformly, one or more R partitions may not fit in memory. Can apply hash-join technique recursively to do the join of this R -partition with corresponding S -partition.



Cost of Hash-Join

- ❖ In partitioning phase, read+write both relns; $2(M+N)$.
In matching phase, read both relns; $M+N$ I/Os.
- ❖ In our running example, this is a total of 4500 I/Os.




Index Nested Loops Join


```
foreach tuple r in R do
    foreach tuple s in S where  $r_i == s_j$  do
        add  $\langle r, s \rangle$  to result
```

- ❖ If there is an index on the join column of one relation (say S), can make it the inner and exploit the index.
 - Cost: $M + (M * p_R) * \text{cost of finding matching S tuples}$
- ❖ For each R tuple, cost of probing S index is about 1.2 for hash index, 2-4 for B+ tree.

General Join Conditions [See Textbook]



- ❖ Equalities over several attributes (e.g., *R.sid=S.sid AND R.rname=S.sname*):
 - For Index NL, build index on *<sid, sname>* (if S is inner); or use existing indexes on *sid* or *sname*.
 - For Sort-Merge and Hash Join, sort/partition on combination of the two join columns.
- ❖ Inequality conditions (e.g., *R.rname < S.sname*):
 - For Index NL, need (clustered!) B+ tree index.
 - Range probes on inner; # matches likely to be much higher than for equality joins.
 - Hash Join, Sort Merge Join not applicable.
 - Block NL quite likely to be the best join method here.



Key Things to Remember

- ❖ Joins are very common, and very expensive
- ❖ Need to bring similar tuples together
- ❖ Can do this in a few ways
 - sort
 - hash
 - using an index
 - or just enumerate all pairs (nested loop joins)